

# Tarefa



## Módulo 7

### Atividade 7.a – Habilitar o log de locks

Ligar parâmetro `log_lock_wait`;

file: `db/data/postgresql.conf`

```
log_lock_waits = on
#log_parameter_max_length = -1

#log_parameter_max_length_on_error = 0

log_statement = 'all'
#log_replication_commands = off
#log_temp_files = -1

# %q = stop here in non-session
#      processes
# %% = '%'
# e.g. '<%u%%d> '
# log lock waits >= deadlock_timeout
# when logging statements, limit logged
# bind-parameter values to N bytes;
# -1 means print in full, 0 disables
# when logging an error, limit logged
# bind-parameter values to N bytes;
# -1 means print in full, 0 disables
# none, ddl, mod, all

# log temporary files equal or larger
```

Recarregar configurações.

### Atividade 7.b. – Testar funcionamento do MVCC

Abrir uma conexão com o banco `benchmark`;

```
postgres@debian10:~$ psql
psql (13.1)
Type "help" for help.

postgres=# \c benchmark;
You are now connected to database "benchmark" as user "postgres".
benchmark=#
```

Iniciar uma transação e fazer um `update` em um registro da tabela `pgbench_accounts`;

```
benchmark=# begin;
BEGIN
benchmark=*# update pgbench_accounts set filler = 'foo' where aid = 10;
UPDATE 1
benchmark=*# select * from pgbench_accounts where aid = 10;
 aid | bid | abalance | filler
-----+-----+-----+-----
  10 |   1 |       0 | foo
(1 row)
```

**Abrir outra conexão e executar um select no mesmo registro;**

```
postgres@debian10:~$ psql -d benchmark
psql (13.1)
Type "help" for help.
```

```
benchmark=# select * from pgbench_accounts where aid = 10;
 aid | bid | abalance |                               filler
-----+-----+-----+-----
  10 |   1 |         0 |
(1 row)
```

**Executar rollback na primeira transação.**

```
benchmark=*# rollback;
ROLLBACK
```

## Atividade 7.c – Testar locks em updates

**Abrir uma conexão com o banco benchmark;**

```
postgres@debian10:~$ psql -d benchmark
psql (13.1)
Type "help" for help.
```

```
benchmark=#
```

**Iniciar uma transação e executar update em um registro;**

```
benchmark=# begin;
BEGIN
benchmark=*# update pgbench_accounts set filler = 'foo' where aid = 10;
UPDATE 1
benchmark=*#
```

**Abrir outra conexão, tentar atualizar o mesmo registro;**

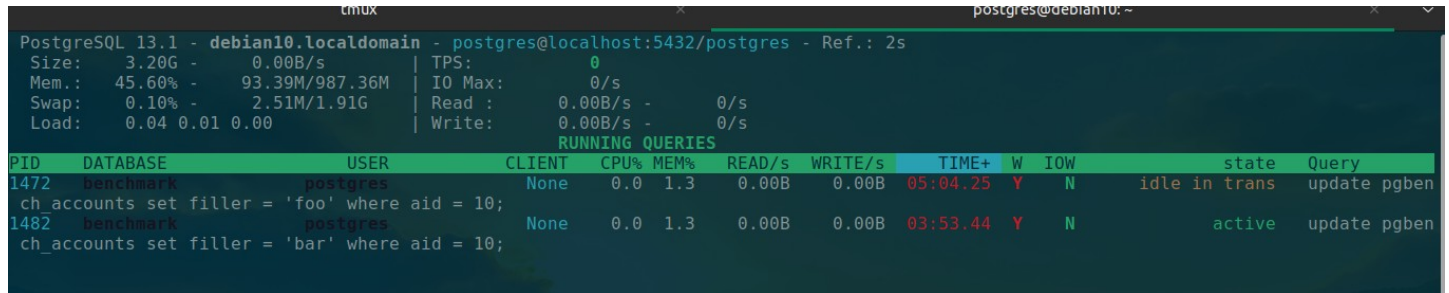
```
postgres@debian10:~$ psql -d benchmark
psql (13.1)
Type "help" for help.
```

```
benchmark=# update pgbench_accounts set filler = 'bar' where aid = 10;
```

**A operação não foi concluída. Ficou bloqueada**

Em um terceiro terminal:

- Verificar locks com pg\_activity;



PostgreSQL 13.1 - debian10.localdomain - postgres@localhost:5432/postgres - Ref.: 2s

PID	DATABASE	USER	CLIENT	CPU%	MEM%	READ/s	WRITE/s	TIME+	W	IOW	state	Query
1472	benchmark	postgres	None	0.0	1.3	0.00B	0.00B	05:04.25	Y	N	idle in trans	update pgben
1482	benchmark	postgres	None	0.0	1.3	0.00B	0.00B	03:53.44	Y	N	active	update pgben

- Localize mensagens de locks na log do PostgreSQL;

```
2022-08-17 19:55:31 UTC [1472]: [3-1] user=postgres,db=benchmark LOG: statement: update
pgbench_accounts set filler = 'foo' where aid = 10;
2022-08-17 19:56:41 UTC [1482]: [1-1] user=postgres,db=benchmark LOG: statement: update
pgbench_accounts set filler = 'bar' where aid = 10;
2022-08-17 19:56:42 UTC [1482]: [2-1] user=postgres,db=benchmark LOG: process 1482 still
waiting for ShareLock on transaction 167101 after 1000.205 ms
2022-08-17 19:56:42 UTC [1482]: [3-1] user=postgres,db=benchmark DETAIL: Process holding
the lock: 1472. Wait queue: 1482.
2022-08-17 19:56:42 UTC [1482]: [4-1] user=postgres,db=benchmark CONTEXT: while updating
tuple (0,10) in relation "pgbench_accounts"
```

- Localize a entrada relacionada no catálogo pg\_locks;

locktype	database	relation	page	tuple	virtualxid	transactionid	classid	objid	objsubid	virtualtransaction	pid	mode	granted	fastpath
relation	12662	12141								5/381	1598	AccessShareLock	t	t
virtualxid					5/381					5/381	1598	ExclusiveLock	t	t
relation	41057	41080								4/79	1482	RowExclusiveLock	t	t
relation	41057	41078								4/79	1482	RowExclusiveLock	t	t
relation	41057	41064								4/79	1482	RowExclusiveLock	t	t
virtualxid					4/79					4/79	1482	ExclusiveLock	t	t
relation	41057	41080								3/138	1472	RowExclusiveLock	t	t
relation	41057	41078								3/138	1472	RowExclusiveLock	t	t
relation	41057	41064								3/138	1472	RowExclusiveLock	t	t
virtualxid					3/138					3/138	1472	ExclusiveLock	t	t
transactionid						167101				4/79	1482	ShareLock	f	f
transactionid						167102				4/79	1482	ExclusiveLock	t	f
tuple										4/79	1482	ExclusiveLock	t	f
transactionid						167101				3/138	1472	ExclusiveLock	t	f

(14 rows)

- Consultar na pg\_stat\_activity as colunas que identificam a situação das conexões envolvidas.

```
postgres=# select state, query from pg_stat_activity;
```

state	query
idle in transaction	update pgbench_accounts set filler = 'foo' where aid = 10;
active	update pgbench_accounts set filler = 'bar' where aid = 10;
active	select state, query from pg_stat_activity;

(8 rows)

Executar um Rollback na primeira transação.

```
benchmark=# rollback;
ROLLBACK
```

## Atividade 7.d – Testar deadlocks

**Abrir uma conexão com o banco benchmark;  
Iniciar uma transação;  
Executar update em um registro;**

```
postgres@debian10:~$ psql -d benchmark
psql (13.1)
Type "help" for help.
```

```
benchmark=# begin;
BEGIN
benchmark=# update pgbench_accounts set filler = 'bar' where aid = 60;
UPDATE 1
benchmark=#
```

**Abrir outra conexão, iniciar uma transação, executar um update em outro registro;**

```
postgres@debian10:~$ psql -d benchmark
psql (13.1)
Type "help" for help.
```

```
benchmark=# begin;
BEGIN
benchmark=# update pgbench_accounts set filler = 'foo' where aid = 50;
UPDATE 1
benchmark=#
```

**Na primeira conexão, tentar atualizar o mesmo registro da segunda;**

```
postgres@debian10:~$ psql -d benchmark
psql (13.1)
Type "help" for help.
```

```
benchmark=# begin;
BEGIN
benchmark=# update pgbench_accounts set filler = 'bar' where aid = 60;
UPDATE 1
benchmark=# update pgbench_accounts set filler = 'barbar' where aid = 50;
```

**Na segunda conexão, tentar atualizar o primeiro registro da primeira conexão.**

```
postgres@debian10:~$ psql -d benchmark
psql (13.1)
Type "help" for help.
```

```
benchmark=# begin;
BEGIN
benchmark=# update pgbench_accounts set filler = 'foo' where aid = 50;
UPDATE 1
benchmark=# update pgbench_accounts set filler = 'foofoo' where aid = 60;
ERROR: deadlock detected
DETAIL: Process 1641 waits for ShareLock on transaction 167105; blocked by process 1637.
Process 1637 waits for ShareLock on transaction 167106; blocked by process 1641.
HINT: See server log for query details.
CONTEXT: while updating tuple (0,60) in relation "pgbench_accounts"
benchmark=#
```

## Atividade 7.e – Usar EXPLAIN e criação de índices

Executar EXPLAIN na seguinte query:

```
SELECT *  
FROM pgbench_accounts a  
INNER JOIN pgbench_branches b ON a.bid=b.bid  
INNER JOIN pgbench_tellers t ON t.bid=b.bid  
WHERE a.bid=56;
```

### Analise o custo

#### QUERY PLAN

```
-----  
Nested Loop (cost=0.43..152264.37 rows=1013070 width=813)  
  -> Index Scan using idx_accounts_bid_aid on pgbench_accounts a (cost=0.43..139578.12  
rows=101307 width=97)  
    Index Cond: (bid = 56)  
    -> Materialize (cost=0.00..22.90 rows=10 width=716)  
      -> Nested Loop (cost=0.00..22.85 rows=10 width=716)  
        -> Seq Scan on pgbench_branches b (cost=0.00..2.25 rows=1 width=364)  
          Filter: (bid = 56)  
        -> Seq Scan on pgbench_tellers t (cost=0.00..20.50 rows=10 width=352)  
          Filter: (bid = 56)  
(9 rows)
```

### Executar EXPLAIN ANALYZE;

- Analise o tempo

#### QUERY PLAN

```
-----  
Nested Loop (cost=0.43..152264.37 rows=1013070 width=813) (actual time=56.642..716.340 rows=1000000 loops=1)  
  -> Index Scan using idx_accounts_bid_aid on pgbench_accounts a (cost=0.43..139578.12 rows=101307 width=97) (actual  
time=40.824..44  
Index Cond: (bid = 56)  
  -> Materialize (cost=0.00..22.90 rows=10 width=716) (actual time=0.000..0.001 rows=10 loops=100000)  
    -> Nested Loop (cost=0.00..22.85 rows=10 width=716) (actual time=15.794..36.233 rows=10 loops=1)  
      -> Seq Scan on pgbench_branches b (cost=0.00..2.25 rows=1 width=364) (actual time=13.509..13.511 rows=1  
loops=1)  
        Filter: (bid = 56)  
        Rows Removed by Filter: 99  
      -> Seq Scan on pgbench_tellers t (cost=0.00..20.50 rows=10 width=352) (actual time=2.277..22.708 rows=10  
loops=1)  
        Filter: (bid = 56)  
        Rows Removed by Filter: 990  
Planning Time: 0.157 ms  
Execution Time: 758.638 ms  
(13 rows)
```

## Executar EXPLAIN (ANALYZE,BUFFERS);

- Analise o acerto em cache e a leitura em disco.

### QUERY PLAN

```
Nested Loop (cost=0.43..152264.37 rows=1013070 width=813) (actual time=0.093..260.140 rows=1000000 loops=1)
  Buffers: shared hit=2027
  -> Index Scan using idx_accounts_bid_aid on pgbench_accounts a (cost=0.43..139578.12 rows=101307 width=97) (actual
time=0.029..21.
    Index Cond: (bid = 56)
    Buffers: shared hit=2018
  -> Materialize (cost=0.00..22.90 rows=10 width=716) (actual time=0.000..0.001 rows=10 loops=100000)
    Buffers: shared hit=9
  -> Nested Loop (cost=0.00..22.85 rows=10 width=716) (actual time=0.059..0.149 rows=10 loops=1)
    Buffers: shared hit=9
    -> Seq Scan on pgbench_branches b (cost=0.00..2.25 rows=1 width=364) (actual time=0.018..0.019 rows=1 loops=1)
      Filter: (bid = 56)
      Rows Removed by Filter: 99
      Buffers: shared hit=1
    -> Seq Scan on pgbench_tellers t (cost=0.00..20.50 rows=10 width=352) (actual time=0.040..0.125 rows=10 loops=1)
      Filter: (bid = 56)
      Rows Removed by Filter: 990
      Buffers: shared hit=8
Planning Time: 0.157 ms
Execution Time: 307.363 ms
(19 rows)
```

## Criar um índice na coluna bid da tabela pgbench\_tellers;

```
benchmark=# CREATE INDEX idx_tellers_bid ON pgbench_tellers(bid);
CREATE INDEX
benchmark=#
```

## Execute um EXPLAIN(ANALYZE,BUFFERS) novamente e analise as informações de custo, tempo e buffers;

### QUERY PLAN

```
Nested Loop (cost=4.79..152256.70 rows=1013070 width=813) (actual time=0.101..255.165 rows=1000000 loops=1)
  Buffers: shared hit=2022 read=2
  -> Index Scan using idx_accounts_bid_aid on pgbench_accounts a (cost=0.43..139578.12 rows=101307 width=97) (actual
time=0.026..21.
    Index Cond: (bid = 56)
    Buffers: shared hit=2018
  -> Materialize (cost=4.35..15.23 rows=10 width=716) (actual time=0.000..0.001 rows=10 loops=100000)
    Buffers: shared hit=4 read=2
  -> Nested Loop (cost=4.35..15.18 rows=10 width=716) (actual time=0.068..0.080 rows=10 loops=1)
    Buffers: shared hit=4 read=2
    -> Seq Scan on pgbench_branches b (cost=0.00..2.25 rows=1 width=364) (actual time=0.017..0.019 rows=1 loops=1)
      Filter: (bid = 56)
      Rows Removed by Filter: 99
      Buffers: shared hit=1
    -> Bitmap Heap Scan on pgbench_tellers t (cost=4.35..12.83 rows=10 width=352) (actual time=0.047..0.055
rows=10 loops=
      Recheck Cond: (bid = 56)
      Heap Blocks: exact=3
      Buffers: shared hit=3 read=2
      -> Bitmap Index Scan on idx_tellers_bid (cost=0.00..4.35 rows=10 width=0) (actual time=0.041..0.041
rows=10 loop
        Index Cond: (bid = 56)
        Buffers: shared read=2
Planning:
  Buffers: shared hit=15 read=1
Planning Time: 0.386 ms
Execution Time: 301.588 ms
(24 rows)
```

## Atividade 7.f – Criar índice composto

Execute o EXPLAIN da query:

```
SELECT DISTINCT *  
FROM pgbench_accounts  
WHERE bid=81 AND aid NOT IN (1,46,28,04,77,93);
```

```
benchmark=# explain  
benchmark=# SELECT DISTINCT *  
benchmark=# FROM pgbench_accounts  
benchmark=# WHERE bid=81 AND aid NOT IN (1,46,28,04,77,93);
```

### QUERY PLAN

```
-----  
HashAggregate (cost=147107.22..148174.14 rows=106692 width=97)  
  Group Key: aid, bid, abalance, filler  
    -> Index Scan using idx_accounts_bid_aid on pgbench_accounts (cost=0.43..146040.30 rows=106692 width=97)  
          Index Cond: (bid = 81)  
          Filter: (aid <> ALL ('{1,46,28,4,77,93}'::integer[]))  
(5 rows)
```

Crie um índice composto para query.

```
benchmark=# CREATE INDEX idx_branch_aid_bid ON pgbench_accounts(aid,bid);  
CREATE INDEX
```

Teste novamente a query e veja o plano de execução

```
benchmark=# explain  
SELECT DISTINCT *  
FROM pgbench_accounts  
WHERE bid=81 AND aid NOT IN (1,46,28,04,77,93);
```

### QUERY PLAN

```
-----  
HashAggregate (cost=146042.34..147099.01 rows=105667 width=97)  
  Group Key: aid, bid, abalance, filler  
    -> Index Scan using idx_accounts_bid_aid on pgbench_accounts (cost=0.43..144985.67 rows=105667 width=97)  
          Index Cond: (bid = 81)  
          Filter: (aid <> ALL ('{1,46,28,4,77,93}'::integer[]))  
(5 rows)
```

## Atividade 7.g – Criar uma Visão Materializada

Abra uma conexão com o banco benchmark;

```
postgres@debian10:~$ psql -d benchmark  
psql (13.1)  
Type "help" for help.
```

```
benchmark=#
```

**Crie uma visão materializada com todos os registros da tabela `pgbench_accounts` cujo `abalance` seja maior que zero;**

NOTA: Caso não haja nenhum registro, rode um teste do `pgbench` primeiro para alterar registros aleatoriamente

```
$ pgbench -T 60 benchmark
```

```
postgres@debian10:~$ psql -d benchmark
psql (13.1)
Type "help" for help.
```

```
benchmark=# create materialized view mv_accounts
benchmark=# as select * from pgbench_accounts
benchmark=# where abalance > 0;
SELECT 2178
benchmark=#
```

## Selecione todos os registros da visão medindo o tempo de execução

Compare com o tempo de execução buscando na tabela original `pgbench_accounts` os registros positivos

DICA: Você pode usar `EXPLAIN ANALYZE` antes da query ou usar `\timing` no `psql`

```
benchmark=# explain (analyze)
benchmark=# select * from mv_accounts;
```

### QUERY PLAN

```
-----
Seq Scan on mv_accounts (cost=0.00..57.78 rows=2178 width=97) (actual
time=0.013..0.533 rows=2178 loops=1)
  Planning Time: 0.229 ms
  Execution Time: 0.825 ms
(3 rows)
```

```
benchmark=# explain (analyze)
select * from pgbench_accounts where abalance > 0;
```

### QUERY PLAN

```
-----
Gather (cost=1000.00..218657.43 rows=1 width=97) (actual time=102.021..19350.573 rows=2178
loops=1)
  Workers Planned: 2
  Workers Launched: 2
  -> Parallel Seq Scan on pgbench_accounts (cost=0.00..217657.33 rows=1 width=97) (actual
time=95.398..19333.292 rows=726 loops=3)
    Filter: (abalance > 0)
    Rows Removed by Filter: 3332607
  Planning Time: 0.173 ms
  Execution Time: 19351.225 ms
(8 rows)
```



## Atividade 7.h - Testar opções de SELECTs com locks

Abrir uma conexão com o banco benchmark;  
Iniciar uma transação;  
Executar update em um registro qualquer; Por exemplo:

```
UPDATE pgbench_accounts SET abalance=100 WHERE aid=1273;
```

```
postgres@debian10:~$ psql -d benchmark
psql (13.1)
Type "help" for help.
```

```
benchmark=# begin;
BEGIN
benchmark=# update pgbench_accounts set filler = 'foo/bar' where aid = 500;
UPDATE 1
benchmark=#
```

Abrir outra conexão, testes os seguintes comandos:

```
SELECT * FROM pgbench_accounts WHERE aid=500 FOR UPDATE NOWAIT;
```

```
benchmark=# SELECT * FROM pgbench_accounts WHERE aid=500 FOR UPDATE NOWAIT;
ERROR:  could not obtain lock on row in relation "pgbench_accounts"
```

```
SELECT * FROM pgbench_accounts WHERE aid=500 FOR UPDATE SKIP LOCKED;
```

```
benchmark=# SELECT * FROM pgbench_accounts WHERE aid=500 FOR UPDATE SKIP LOCKED;
 aid | bid | abalance | filler
-----+-----+-----+-----
(0 rows)
```

```
SELECT * FROM pgbench_accounts WHERE aid=500 FOR UPDATE;
```

```
benchmark=# SELECT * FROM pgbench_accounts WHERE aid=500 FOR UPDATE;
```

PostgreSQL 13.1 - debian10.localdomain - postgres@localhost:5432/postgres - Ref.: 2s									
Size:	3.41G	-	0.00B/s		TPS:	0			
Mem.:	44.40%	-	82.55M/987.36M		IO Max:	5/s			
Swap:	0.70%	-	14.00M/1.91G		Read :	0.00B/s	-	0/s	
Load:	0.13	0.09	0.09		Write:	0.00B/s	-	0/s	
WAITING QUERIES									
PID	DATABASE	RELATION	TYPE	MODE	TIME+	state	Query		
1754	benchmark	None	transactionid	ShareLock	01:44.52	active	SELECT * FROM pgbench_accounts WHERE aid=500 FOR UPDATE;		

### Interprete as diferenças entre os comandos

A opção **for update** faz com que a consulta fique bloqueada até o registro ser liberado. **Nowait**, gera um erro, mas não bloqueia a consulta. **Skipe locked** ignora os registros bloqueados.